

Lecture 6: CS677

Sept 7, 2017

Review

- HW1 due September 12
- Cloud computing: students can get a better personal account at:
https://console.cloud.google.com/freetrial?_ga=2.228461851.-722665125.1503520492&page=1
- Previous class
 - Shape from shading and photometric stereo
 - Color perception
 - OpenCV tutorial
- Today's objective
 - Image filtering
 - Image segmentation

Comments on HW1

- Final formulas for vanishing points and lines are given in a slide in Lec4 notes
- These formulas were not derived, just stated
- For HW1, students need to derive these formulas from the basic projection equations
- Answers need not be expressed in the given form, nor derived using concepts of plane or line at infinity; other ways of approaching the solution are equally acceptable.

Next Topics

- We are finished with image formation
- Next major topics of interest
 - Inference of 3-D
 - Single Images (difficult)
 - From multiple views
 - From direct range sensing
 - Detection and recognition of objects
- Image feature extraction and segmentation
 - Useful for both of the above topics
 - This is what we will study next


Image Smoothing/Filtering

- Reduce effects of noise, texture
- Enhance sharpness of edges
- Linear Filters
- General process
 - Form a new image whose pixels are a weighted sum of original pixel values, using the same set of weights at each point
 - Smoothing by averaging
 - Form the average of pixels in a neighborhood

$$\mathcal{R}_{ij} = \frac{1}{(2k+1)^2} \sum_{u=i-k}^{u=i+k} \sum_{v=j-k}^{v=j+k} \mathcal{F}_{uv}.$$

- Weights need not be uniform: high value at center, drops off as a function of distance (e.g. as a Gaussian function)
- An example on next slide

Example: Smoothing



Original Image

Average smoothing

Gaussian Smoothing

Convolution

- Represent weights as an image, H (sometimes called the **kernel**)
- *Convolving* a filter H with image F gives a new image R given by:

$$R_{ij} = \sum_{u,v} H_{i-u,j-v} F_{uv}$$

- Note: reversal of sign of dummy variables, u and v ; sum is over all non-zero values
- Convolution commonly written as $R = H * F$
- Convolution formulation provides many useful properties studied in basic signal processing, we list only a few here
- Convolution is linear and shift invariant
 - Scaling of H or F results in linear scaling of R
 - Shift invariance: translation of signal results in same translation of result
 - Commutative: $g*h = h*g$; associative $(f*(g*h)) = (f*g)*h$
- Edge effects
 - Values undefined at image boundary: pad the image by zeros beyond the border

Derivative (Difference) by Convolution

- Consider the convolution filter given below:

$$\mathcal{H} = \begin{Bmatrix} 0 & 0 & 0 \\ 1 & 0 & -1 \\ 0 & 0 & 0 \end{Bmatrix}$$

- Computes difference in the x-direction
- Following compute derivatives in x and y directions with some influence from neighboring rows and columns (Sobel masks)

| | | | | | |
|----|---|---|----|----|----|
| -1 | 0 | 1 | 1 | 2 | 1 |
| -2 | 0 | 2 | 0 | 0 | 0 |
| -1 | 0 | 1 | -1 | -2 | -1 |




Image Pyramids

- Interesting patterns may occur at different scales
 - e.g. on a zebra, are we interested in hair, stripes or just the neck?
- Applying large filters is expensive
 - We can get same effect by reducing resolution of the image
- Figure shows multiple levels of reduction, resulting in a *pyramid*
 - We can sample alternate pixels (for example) or convolve with a Gaussian of selected σ .


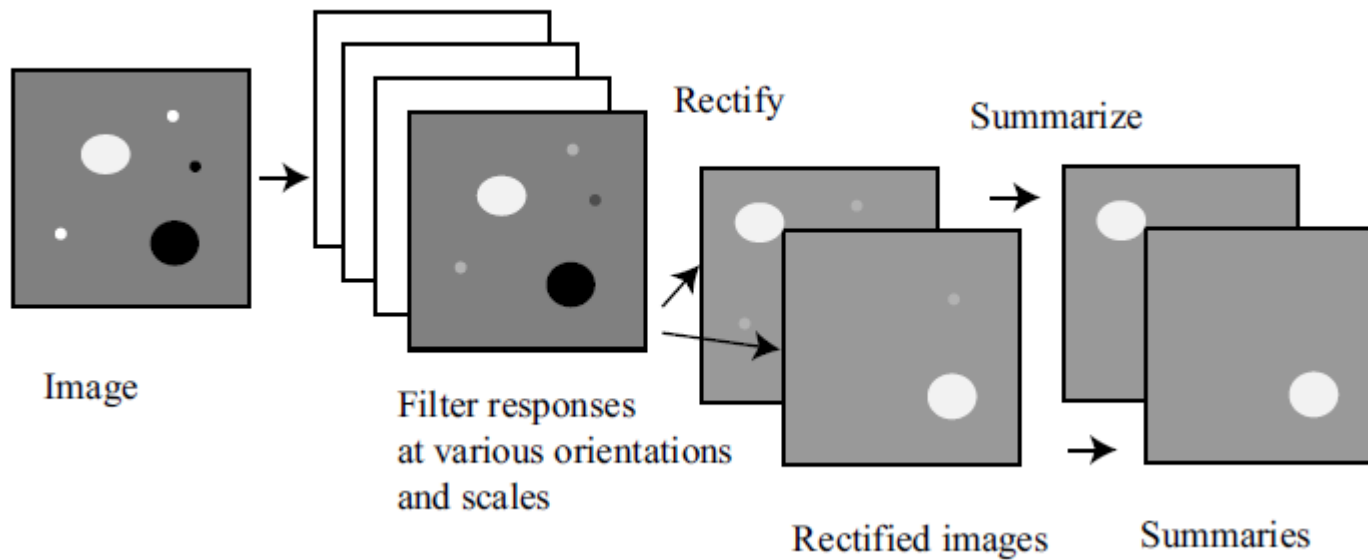


Texture

- Different materials may have different characteristic patterns, that we call texture



Texture Representation



Filter bank: oriented and isotropic

Texture Comments

- Texture seems characterized by repeated patterns
 - However, natural textures are not completely uniform
- What is the size of image patch over which we should compute texture patterns?
- For natural textures, one approach is to compute outputs of various filters
 - Outputs of filters that emphasize different patterns, at different scales and different orientations
- Texture variation (gradient) can be a strong indicator of surface orientation
 - But requires texture to be homogeneous

Image Segmentation


- Find boundaries of objects and surfaces in the scene
 - Typically characterized by discontinuities in range/depth of points in the scene (assumes object surfaces are continuous)
 - However, depth information is not readily available; instead we detect intensity (or color) discontinuities which may not always correspond to object boundaries
- Some examples on following slides

Image Segmentation



Example from Berkeley Segmentation Dataset

Image Segmentation



Example from Berkeley Segmentation Dataset

Image Segmentation: Approaches

- Find points of rapid changes in intensity/color (edge detection), connect them to give boundaries
- Find *regions* of constant properties (intensity, color, texture..)
- We will study region segmentation first (chapter 9); book does it in the reverse order.

Region Segmentation (FP 9.3,9.4)

- Segmentation is equivalent to *clustering*
 - ***Cluster***: collection of data samples (pixels)
- Divisive
 - Start with entire data set as a single cluster
 - Split cluster into components that yield largest inter-cluster distance
 - Repeat until distance or clusters become small
- Agglomerative
 - Start with each point being a separate cluster
 - Merge clusters with smallest inter-cluster distance
 - Repeat until distance or clusters become large

K-Means Clustering

- Choose a fixed number of clusters
- Choose cluster centers and point-cluster allocations to minimize error:

$$\sum_{i \in \text{clusters}} \left\{ \sum_{j \in \text{elements of } i\text{'th cluster}} \|x_j - \mu_i\|^2 \right\}$$

- Too many possible allocations for exhaustive search
- Algorithm
 - Fix cluster centers; allocate points to closest cluster
 - Fix allocation; compute the cluster centers
 - Initial assignment may be random
- x can be any *vector* in a space for which we can compute a distance

Image




Clusters on intensity



Clusters on color



K-means clustering using intensity only and using color
(five clusters)



K-means using color,
11 segments.


Left: image with mean values
of clusters

Others: Four segments shown,
note a segment is not
necessarily connected



Connected Components

- Compute *connected components*
 - 4 or 8 connectivity
 - Connectivity algorithms not covered (but are simple)




| | | | | | | | |
|--|--|---|---|---|---|--|--|
| | | | | | | | |
| | | | | | | | |
| | | | 1 | 1 | | | |
| | | 1 | | | 1 | | |
| | | 1 | | | 1 | | |
| | | | 1 | 1 | | | |
| | | | | | | | |
| | | | | | | | |

Superpixels

- Goal is to produce segments of similar size with high fidelity to boundaries
- Segments not likely to correspond to objects; objects to be detected in a subsequent stage
- May be useful for various forms of post-processing
- Has become a popular first step in recent years
- Watershed algorithm
- SLIC algorithm (not in book)

Watershed Algorithm

- Imagine intensity of image to represent terrain height
- Find “catchment basins” (lakes) that would form from rain fall
- Find local minima: consider them to be seeds and give a unique label to each
- For any pixel, go in direction of negative gradient, if a seed is reached (or a labeled pixel is reached), take its label.
- Efficient algorithms $O(N \cdot \log N)$ can be constructed



On image intensity

On gradient magnitude

Next Class

- Chapter 9, section 9.4
- Chapter 5, sections 5.1, 5.2