

Lecture 8: CS677

Sept 14, 2017

Admin

- HW2 to be posted today, due Sept 24, 9AM
- Exam 1, Oct 10, class period
 - Closed book, closed notes
 - Topics: whatever we cover until Oct 5
 - A detailed list will be provided later
- Make-up classes
 - Instructor needs to be absent Oct 24 and 26 (ICCV)
 - As our classes are recorded, we can make up by holding extra sessions on Oct 13 and Oct 20 (both Fridays)
 - Physical attendance is not required
 - Available slots are 8-9:50AM or 3:30-5:00 pm
 - Latter slot can accommodate only 24 students in class
 - Which do we prefer?

Review

- Previous class
 - Superpixel algorithm: SLIC
 - Mean-shift algorithm
 - Graph-based methods: Normalized cuts
 - Note: both mean-shift and normalized cut papers cited >10K times
- Today's objective
 - Two more graph-based algorithms
 - Edge and curve detection

FH Region Growing Method

- Agglomerative method, FP 9.4.2
- Algorithm does not have a memorable name, so referred to just by author names (Felzenswalb and Huttenlocher, we will call it F-H method).
- Construct a graph as in normalized cuts. Let $w(v_1, v_2)$ be weight of edges connecting vertices v_1 and v_2 . Weight is high if pixels are *dissimilar* (opposite of the convention in normalized cuts)
- Start from small clusters, merge as appropriate
- Informally, the idea is to merge clusters if weight of best edge connecting them is less than the internal differences of the two clusters.
- Formal definitions on next page

FH Region Growing Method

- Difference between two components is defined by:

$$\text{diff}(\mathcal{C}_1, \mathcal{C}_2) = \min_{v_1 \in \mathcal{C}_1, v_2 \in \mathcal{C}_2, (v_1, v_2) \in \mathcal{E}} w(v_1, v_2)$$

- To define internal difference of a component, \mathcal{C} , first construct a minimum spanning tree, $M(\mathcal{C})$, then

$$\text{int}(\mathcal{C}) = \max_{e \in M(\mathcal{C})} w(e)$$

- Now define: $\text{MInt}(\mathcal{C}_1, \mathcal{C}_2) = \min(\text{int}(\mathcal{C}_1) + \tau(\mathcal{C}_1), \text{int}(\mathcal{C}_2) + \tau(\mathcal{C}_2))$

$$\tau(\mathcal{C}) = k / |\mathcal{C}| \quad k \text{ is a constant parameter}$$

$\tau(\mathcal{C})$ creates a bias against small clusters

- Merge \mathcal{C}_1 and \mathcal{C}_2 if $\text{diff}(\mathcal{C}_1, \mathcal{C}_2) < \text{MInt}(\mathcal{C}_1, \mathcal{C}_2)$
- Start with one cluster per pixel, merge clusters iteratively
- Sort edges by weight, start with the smallest until a merge happens
 - Stop if no changes take place
- Method much faster than normalized cuts but less popular (“only” 4K citations), perhaps because it is more complex to implement.
- Not in OpenCV but code available from the authors

FH Results



Energy Minimization Approach

- FP 9.4.3 but following is presented a bit differently
- General Idea
 - Task is to label each pixel as being *foreground* or *background*
 - Requires a model of FG/BG (*e.g.* an intensity/color *pdf*)
 - Label assigned to a pixel depends not only on its local properties but also on the labels of the pixels in the neighborhood, to enforce some smoothness.
- Commonly expressed as the task of minimizing an energy function

$$E(f) = E_{smooth}(f) + E_{data}(f)$$

- Data and smoothness functions need to be defined by the designer
 - For example, data energy can be the absolute difference between observed intensity and expected intensity (based on FG/BG models)
 - Smoothness term on next slide

Energy Minimization Approach

- Smoothness energy can be defined as being sum of terms for each neighbor considered
 - Common to use four neighbors
 - One commonly used smooth energy term:
 - The term is zero if the two neighbors have the same label, and another value if they are different.
 - This value may be different for different labels (background or foreground) and for different neighbors.
- For certain classes of energy functions (“sub-modular” functions), the total energy can be minimized by solving a related graph-cut problem (different from normalized cut approach)
 - See next slide

Graph Cut Formulation

- S and T are two “terminal” nodes corresponding to object and background
- Each pixel node has links to both S and T nodes: strength (capacity) depends on similarity of node to background/foreground
- Links between neighboring nodes; strength depends on similarity of labels
- Minimum energy (max probability) found by min-cut or max-flow algorithms
- We omit details of how to compute the min-cut

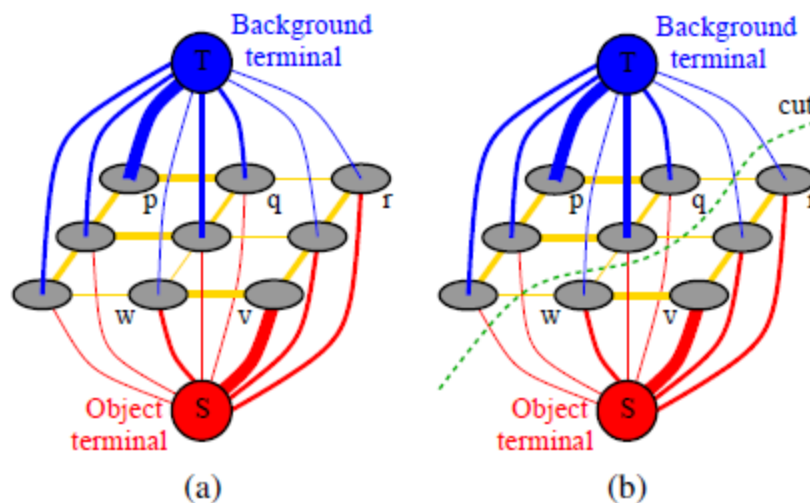


Fig 5.23, RS Book

Grabcut

- A bounding box is provided around the object of interest.
- This box is used to compute properties (such as a color distribution) of the object and the background
- Min-cut method is used to segment
- After initial segmentation, improved models can be obtained and the process can be repeated iteratively.
- Further interactions may be used for further refinement.
- Some examples on following slides.

Grabcut Example



(a)



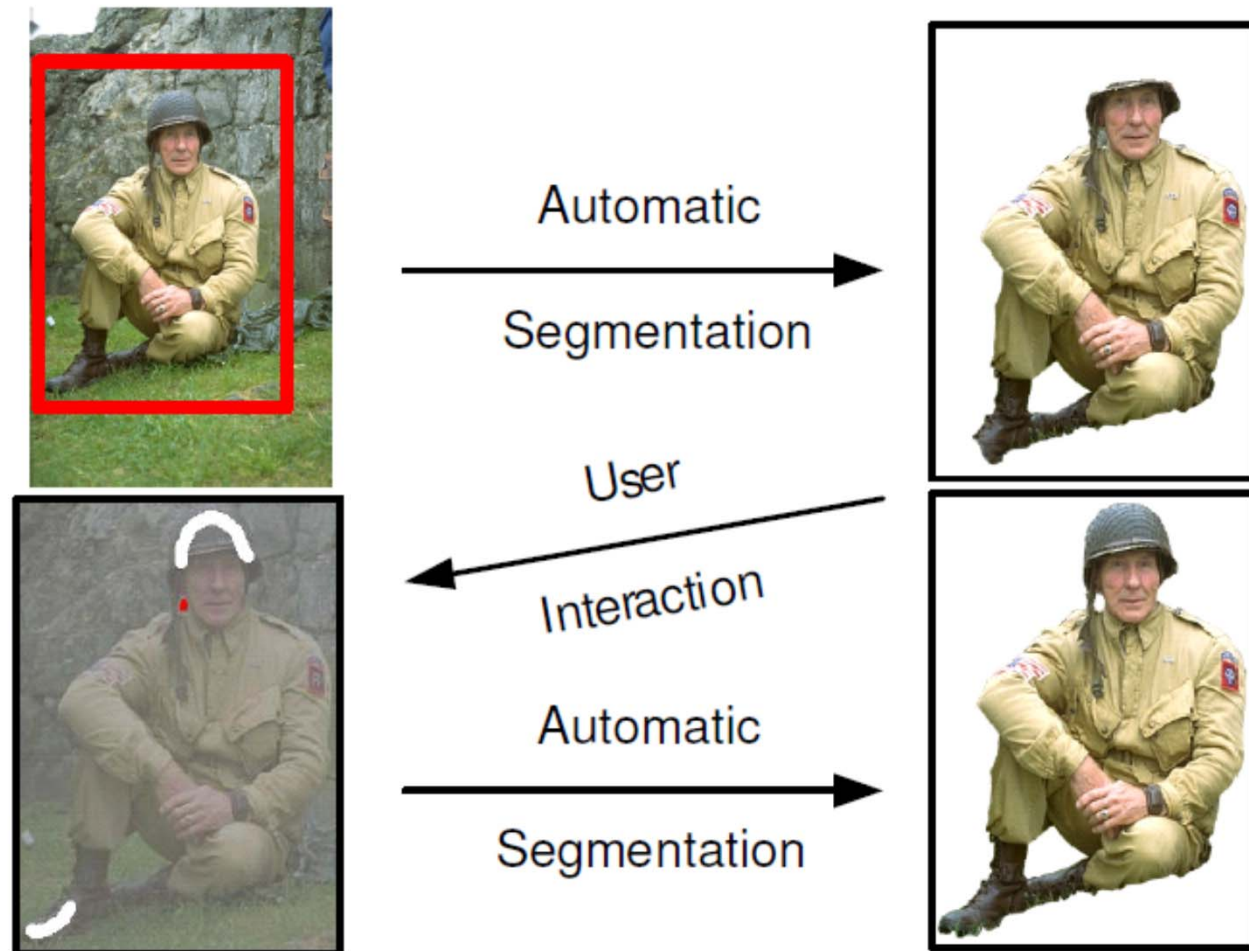
(b)



(c)

Figure 5.24: *GrabCut image segmentation (Rother et al. 2004) © 2004 ACM: (a) the user draws a bounding box in red; (b) the algorithm guesses color distributions for the object and background and performs a binary segmentation; (c) the process is repeated with better region statistics.*

Another Example with Editing



Region Segmentation Summary

- Several divisive and agglomerative methods
- Performance depends on choice of measures of uniformity (affinity), energy functions, thresholds in some cases
- Common problems
 - Regions of interests are not necessarily uniform in a property
 - Shape and context affect segmentation but hard to accommodate in current computer algorithms.

Edge Detection

- Compute discontinuities in image properties (intensity, color, texture)
 - More local approach than region segmentation
 - Better localization, incomplete boundaries, *over*-segmentation
- Approach
 - Measurement of derivatives (Gradient, Laplacian.....)

Computing Derivatives

- 1-D function, $f(x)$
 - Differences of neighboring values
 - $\Delta(x) = f(x) - f(x+1)$ or $f(x) - f(x-1)$ or $f(x+1) - f(x-1)$
 - Consider more neighbors, weight the differences
- 2-D function, $f(x,y)$
 - Can compute derivatives *w.r.t.* x and y separately, say f_x, f_y
 - Gradient is the vector (f_x, f_y)
 - Magnitude: $\sqrt{f_x^2 + f_y^2}$
 - Angle: $\tan^{-1}(f_y / f_x)$
 - Laplacian (non-directional) :

$$(\nabla^2 f)(x, y) = \frac{\partial^2 f}{\partial x^2} + \frac{\partial^2 f}{\partial y^2}$$

Convolution

- Computing derivatives and gradients is equivalent to convolution of image with an *edge mask*
- Examples (1-D: difference, 2-D: Sobel edge mask)

$$\begin{array}{c|c|c} -1 & 0 & 1 \end{array}$$

$$\begin{array}{c|c|c} -1 & 0 & 1 \\ \hline -2 & 0 & 2 \\ \hline -1 & 0 & 1 \end{array} \quad \begin{array}{c|c|c} 1 & 2 & 1 \\ \hline 0 & 0 & 0 \\ \hline -1 & -2 & -1 \end{array}$$

Smoothing and Computing Gradients

- Common to smooth with Gaussian (scale parameter to be supplied) prior to differentiation
- Convolve image with Gaussian and then differentiate
- Equivalent to differentiate the Gaussian and then convolve

$$\frac{\partial (G_{\sigma} ** I)}{\partial x} = \left(\frac{\partial G_{\sigma}}{\partial x} \right) ** I$$

- Notation: “**” stands for convolution
- Derivative of smoothing functions can be pre-computed

Canny Edge Detector

- >27000 citations
- Smooth with Gaussian (need to supply a scale parameter)
- Compute gradient direction
- Compute (estimate) derivative in the direction of the gradient
 - Non-maxima suppression to detect position
- Thresholding
 - High threshold: always accept points with higher values
 - Low threshold: always reject points with lower values
 - Accept points with in between values if connected to an existing edge (hysteresis)
- Some details may be found in:

http://www.dai.ed.ac.uk/CVonline/LOCAL_COPIES/MARBLE/low/edges/canny.htm

Canny Examples

- Three scales




Image of girl



$\sigma=1, m=9, l=0.3, h=0.9$



$\sigma=0.5, m=7, l=0.3, h=0.9$



$\sigma=1, m=9, l=0.1, h=0.3$

- How to choose the “right” scale
 - From *a priori* knowledge of desired objects
 - Use large sigma to find prominent edges, reduce smoothing for better localization, requires tracking across scales

Color Edges

- Color is a 3-D quantity
 - How to compute its gradient?
 - Compute edges in each component and combine (use OR or AND operator?)
 - Combine three components in one (basically gives intensity)
 - Computing derivative of hue is difficult (angle is measured *modulo* 360 degrees)
- Some Observations
 - Intensity and hue edges are highly correlated
 - Unlikely that two surfaces have different hue but same intensity
 - Humans do not localize pure hue edges well

Edgels to Curves

- Edge detectors give edge points or *edgels*
- Curve is defined by an ordered list of points
 - Connected components (regions) do not necessarily contain order information
- *Linking*
 - Given an edge, find the next edge
 - Look for a neighbor in the direction of the edge (normal to gradient)
 - May impose constraints on continuity of edgel direction
 - If multiple choices available, select one
 - Based on local measures (direction/contrast continuity)
 - Global search (corresponds to finding best path in a connectivity graph)

Representation of a Curve

- List of points (in order)
- Fit lines (or some other curves) to the list of points
 - Curve may not be well approximated by a single line
 - Segment at “corners”
- Iterative end point fit
 - Join first and last point by straight line
 - Compute maximum distance of points on the chain, segment at this point if error $>$ threshold, Iterate
 - Example
 - Dual: incremental line fit, keep adding points until error $> t$ (FP Algorithm 10.1)
 - Note: corners may not be where humans would put them but the fit will be within error threshold

Fitting Curves

- Fit curves (straight lines, quadratics, polynomials..) to segmented set of points
 - Least mean squares fit (FP 10.2.1)
 - Yields a compact representation
 - However, adjacent curves may not be continuous
- Fit *splines*
 - Piecewise polynomials, preserve $(n-1)$ order continuity at junctions (knots) for polynomial of order n
 - Cubic splines keep 2nd order derivatives the same
 - Segmentation may be part of fitting process
 - Details of spline fit not covered
 - Normal topic in numerical analysis courses
 - Methods usually available in numerical software libraries

Next Class

- Hough Transform, FP 10.1
- SIFT FP 5.4.1, tutorial to be posted separately